Stage-by-Stage Impact of Blind Hiring

1 Treatment Effect on Evaluations

1.1 On All Dimensions (Z)

Takeaway: Blinding demographic information *reshuffles* candidates' relative evaluations. In the HR resume screening stage, Hispanic Male candidates are the largest losers: under blinding, their overall evaluation scores fall by 11.5% of a standard deviation relative to when their identity is revealed. In the Engineer technical evaluation stage, the largest gainers are White Male candidates, whose scores rise by 8.5% of a SD under blinding, while Hispanic Females' scores lose by 6.9% of a SD. However, because hiring is often an ordinal ranking problem — where one candidate's gain necessarily corresponds to another candidate's lost opportunity — the absence of statistical significance on a given row does not mean the corresponding group is unaffected by blind hiring.

Panel A: HR

	Coding (Z)	Fit (Z)	Stay (Z)	Overall (Z)	Advance (0/1)
White Male	0.038	0.073	0.072	0.037	−0.006
	(0.056)	(0.057)	(0.059)	(0.057)	(0.050)
White Female	0.020	0.010	0.022	0.013	0.003
	(0.046)	(0.053)	(0.065)	(0.051)	(0.049)
Asian Male	0.032	0.029	−0.050	0.015	−0.022
	(0.052)	(0.061)	(0.061)	(0.059)	(0.049)
Asian Female	0.020	−0.011	0.028	−0.023	−0.019
	(0.052)	(0.056)	(0.062)	(0.048)	(0.052)
Hispanic Male	−0.142***	−0.056	−0.084	−0.115**	−0.047
	(0.050)	(0.053)	(0.062)	(0.048)	(0.049)
Hispanic Female	0.040	−0.049	0.013	0.087*	−0.001
	(0.058)	(0.067)	(0.074)	(0.052)	(0.056)
Observations	2466	2430	2394	2466	2466
R2	0.398	0.315	0.196	0.423	0.229
Resume (Content) FE	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓
Respondent FE	×	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓	✓

* p < 0.1, ** p < 0.05, *** p < 0.01

Panel B: Engineer

	Coding (Z)	Fit (Z)	Stay (Z)	Overall (Z)
White Male	0.072**	0.040	0.045	0.085***
	(0.029)	(0.041)	(0.049)	(0.028)
White Female	0.028	0.046	0.023	0.023
	(0.031)	(0.045)	(0.053)	(0.036)
Asian Male	−0.071**	0.083*	0.002	−0.018
	(0.028)	(0.046)	(0.050)	(0.028)
Asian Female	0.003	−0.099**	−0.053	−0.029
	(0.029)	(0.042)	(0.048)	(0.029)
Hispanic Male	0.009	0.023	0.042	−0.003
	(0.027)	(0.037)	(0.054)	(0.030)
Hispanic Female	−0.050*	−0.101**	−0.067	−0.069**
	(0.027)	(0.042)	(0.059)	(0.034)
Observations	2772	2772	2736	2772
R2	0.788	0.571	0.345	0.743
Resume (Content) FE	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓
Respondent FE	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓

* p < 0.1, ** p < 0.05, *** p < 0.01

Standard errors clustered by respondent. Note on Interpreting Coefficients: The underlying regression uses Blind resumes as the reference category, so raw coefficients capture (nonblind − blind) or the effect of being blinded for each demographic on standardized overall evaluation scores. We then multiply by −1 to report (blind − nonblind), aligning the sign convention with the standard "effect of blinding" framing. positive coefficient ⇒ blinding raises the score for that demographic.

1.2 Does Blinding Lead to More Accurate Coding Evaluations?

Takeaway: Evaluators do extract signal about candidates' actual coding ability — Engineers much more strongly than HR, consistent with Engineers seeing the code while HR sees only the resume. Blinding, however, does *not* meaningfully change this signal-extraction ability in either role: the `true_coding_z` $\times$ `blind` interaction is small, negative, and statistically insignificant for both HR and Engineers. If blinding — which in our experiment amounts to removing random, uninformative demographic noise — improved signal-extraction ability, we would expect a positive and statistically significant interaction term. In short, blinding *reshuffles* who receives high scores (Section 1.1) but does not change *how well* those scores track actual coding skill.

Specification: For each role $\in \{\text{HR}, \text{Engineer}\}$, we fit

$$\text{Hiring Evaluation Score (Z)} \sim \text{True Coding Ability (Z)} \times \mathbb{1}[\text{Blind arm}] \mid \text{Respondent FE} + \text{Display Position FE}.$$

In each of the four cells we test two nulls:

- $H_0: \beta = 0$ — the evaluator extracts *no signal* from true coding ability.
- $H_0: \beta = 1$ — the evaluator is *perfectly calibrated*: a 1 SD rise in true ability produces exactly a 1 SD rise in the rating, with no shrinkage.

Running both tests separately by arm reveals whether blinding changes the evaluator's calibration, not merely whether signal is extracted at all.

	HR	Engineer
$\hat{\beta}_{\text{nonblind}}$ (slope, nonblind arm)	0.377 (0.029)	0.886 (0.012)
$\hat{\beta}_{\text{blind}}$ (slope, blind arm)	0.357 (0.023)	0.856 (0.015)
$H_0: \beta_{\text{nonblind}} = 0$	p = <0.001***	p = <0.001***
$H_0: \beta_{\text{nonblind}} = 1$	p = <0.001***	p = <0.001***
$H_0: \beta_{\text{blind}} = 0$	p = <0.001***	p = <0.001***
$H_0: \beta_{\text{blind}} = 1$	p = <0.001***	p = <0.001***
$H_0: \beta_{\text{blind}} = \beta_{\text{nonblind}}$	p = 0.583	p = 0.125
Observations	2,466	2,772
R^2	0.142	0.759
Adj. R^2	0.085	0.743
Resume (Content) FE	×	×
Display Position FE	✓	✓
Respondent FE	✓	✓
Respondent Characteristics	×	×

Cluster-robust standard errors (by respondent) shown in parentheses next to each $\hat{\beta}$. Stars on p-values: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. *Note:* Both arm-specific slopes come from a single pooled interacted regression per role, `sc_coding_z ~ true_coding_z * i(treat, ref = 'nonblind') | responseid + r_do`; $\hat{\beta}_{\text{nonblind}}$ is read off the main coefficient, and $\hat{\beta}_{\text{blind}}$ is computed as (main + interaction) with SE derived via the standard linear-combination formula using the estimated variance–covariance matrix. The "Resume (Content) FE" row is marked as not included because `overall_score` is a resume-level property (one lab score per resume), so a resume-content FE would absorb the RHS and leave the main slope unidentified. "Respondent Characteristics" is marked as not included because respondent-level variables (age, gender, race, education, income) are all perfectly collinear with the respondent FE and therefore absorbed by it.

2 Treatment Effect: Heterogeneity by Demographics

Takeaway: For HR, the effect of blinding does not vary with respondent–applicant demographic concordance. For Engineers, heterogeneity in the treatment effect emerges as the concordance definition tightens. On detailed race–gender, the finest cut of own-group similarity, blinding raises concordant applicants' Overall scores while reducing the nonconcordant group's scores. Concretely: candidates with concordant race and gender profiles gain 7.1% of a SD, and those with nonconcordant profiles lose 1.1% of a SD, on their evaluation scores under blinding. This is consistent with the earlier finding in `8A_Preferences.pdf` §2 that Engineers in our sample penalize candidates who match them on race and gender — a pattern driven mainly by White Male Engineers' assignment of a premium to non-White candidates.

2.1 By Concordant Race

	Panel A: HR (N = 137)					Panel B: Engineer (N = 154)			
	Coding	Fit	Stay	Overall	Advance	Coding	Fit	Stay	Overall
Effect of Blinding on Concordant Applicant	−0.011	−0.001	−0.041	−0.020	−0.042	0.032	0.046	0.025	0.045*
	(0.039)	(0.040)	(0.047)	(0.042)	(0.046)	(0.021)	(0.031)	(0.039)	(0.023)
Effect of Blinding on Nonconcordant Applicant	0.004	0.000	0.014	0.007	−0.007	−0.010	−0.014	−0.008	−0.014*
	(0.014)	(0.014)	(0.017)	(0.015)	(0.044)	(0.006)	(0.009)	(0.012)	(0.007)
Observations	2466	2430	2394	2466	2466	2772	2772	2736	2772
R^2	0.396	0.314	0.195	0.421	0.229	0.787	0.569	0.344	0.742
Adj. R^2	0.382	0.297	0.176	0.407	0.211	0.783	0.560	0.331	0.737
Resume (Content) FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Respondent FE	×	×	×	×	×	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓	✓	✓	✓	✓	✓

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Standard errors clustered by respondent.

Note: Concordance is defined at the detailed race level: Asian respondents are matched at the sub-ethnicity level (Chinese / Indian). Respondents whose race is not represented in the applicant pool (e.g., Black, Other; Korean / Japanese / other-Asian whose sub-ethnicity is not Chinese or Indian) are treated as always-discordant in nonblind.

Coefficients are flipped so positive ⇒ blinding raises the score for that concordance group relative to the nonblind outcome.

2.2 By Concordant Gender

	Panel A: HR (N = 137)					Panel B: Engineer (N = 154)			
	Coding	Fit	Stay	Overall	Advance	Coding	Fit	Stay	Overall
Effect of Blinding on Concordant Applicant	−0.018	−0.024	0.007	−0.010	−0.022	0.005	0.033*	0.036	0.029*
	(0.026)	(0.029)	(0.035)	(0.027)	(0.047)	(0.012)	(0.019)	(0.028)	(0.015)
Effect of Blinding on Nonconcordant Applicant	0.017	0.021	−0.007	0.009	−0.010	−0.005	−0.034*	−0.037	−0.030*
	(0.024)	(0.026)	(0.031)	(0.025)	(0.043)	(0.013)	(0.019)	(0.029)	(0.016)
Observations	2466	2430	2394	2466	2466	2772	2772	2736	2772
R^2	0.396	0.314	0.195	0.421	0.229	0.787	0.569	0.345	0.742
Adj. R^2	0.382	0.298	0.176	0.407	0.211	0.783	0.560	0.331	0.737
Resume (Content) FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Respondent FE	×	×	×	×	×	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓	✓	✓	✓	✓	✓

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Standard errors clustered by respondent.

Note: Respondents whose gender is not represented in the applicant pool (e.g., Non-binary, Other) are treated as always-discordant in nonblind, since the applicant names are exclusively Male or Female.

Coefficients are flipped so positive ⇒ blinding raises the score for that concordance group relative to the nonblind outcome.

2.3 By Concordant Race and Gender

	Panel A: HR (N = 137)					Panel B: Engineer (N = 154)			
	Coding	Fit	Stay	Overall	Advance	Coding	Fit	Stay	Overall
Effect of Blinding on Concordant Applicant	−0.030	0.008	0.018	−0.057	−0.022	0.004	0.003	−0.011	0.040
	(0.058)	(0.065)	(0.078)	(0.063)	(0.058)	(0.032)	(0.044)	(0.050)	(0.028)
Effect of Blinding on Nonconcordant Applicant	0.005	−0.001	−0.003	0.010	−0.015	−0.001	−0.001	0.002	−0.008
	(0.010)	(0.012)	(0.014)	(0.011)	(0.043)	(0.007)	(0.009)	(0.010)	(0.006)
Observations	2466	2430	2394	2466	2466	2772	2772	2736	2772
R^2	0.396	0.314	0.195	0.421	0.228	0.787	0.569	0.344	0.742
Adj. R^2	0.382	0.297	0.176	0.408	0.211	0.783	0.560	0.331	0.737
Resume (Content) FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Respondent FE	×	×	×	×	×	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓	✓	✓	✓	✓	✓

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Standard errors clustered by respondent.

Note: Concordance is defined as race-gender match at the broad level (Asian respondents are matched at the broad Asian level rather than by detailed sub-ethnicity). Respondents whose race or gender is not represented in the applicant pool (e.g., Black, Other; Non-binary, Other-gender) are treated as always-discordant in nonblind.

Coefficients are flipped so positive ⇒ blinding raises the score for that concordance group relative to the nonblind outcome.

2.4 By Concordant Race and Gender (using detailed Asian categorization)

	Panel A: HR (N = 137)					Panel B: Engineer (N = 154)			
	Coding	Fit	Stay	Overall	Advance	Coding	Fit	Stay	Overall
Effect of Blinding on Concordant Applicant	−0.030	0.020	0.019	−0.076	−0.029	0.037	0.047	0.049	0.069**
	(0.064)	(0.067)	(0.087)	(0.068)	(0.059)	(0.029)	(0.047)	(0.060)	(0.032)
Effect of Blinding on Nonconcordant Applicant	0.004	−0.003	−0.003	0.011	−0.014	−0.005	−0.007	−0.007	−0.009**
	(0.009)	(0.010)	(0.012)	(0.010)	(0.043)	(0.004)	(0.006)	(0.008)	(0.004)
Observations	2466	2430	2394	2466	2466	2772	2772	2736	2772
R^2	0.396	0.314	0.195	0.421	0.229	0.787	0.569	0.344	0.742
Adj. R^2	0.382	0.297	0.176	0.408	0.211	0.783	0.560	0.331	0.737
Resume (Content) FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓	✓	✓	✓	✓
Respondent FE	×	×	×	×	×	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓	✓	✓	✓	✓	✓

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Standard errors clustered by respondent.

Note: Concordance is defined at the detailed race-gender level: applicants are Chinese, Indian, White, or Hispanic (by gender). Respondents whose detailed race-gender is not represented in the applicant pool (e.g., Korean / Japanese / other-Asian sub-ethnicities; Black, Other-race; Non-binary, Other-gender) are treated as always-discordant in nonblind.

Coefficients are flipped so positive ⇒ blinding raises the score for that concordance group relative to the nonblind outcome.

3 Treatment Effect: Heterogeneity by DEI / Diversity Preferences

Takeaway: [To be filled after first render. Parallel to 8A §3, we ask whether the effect of blinding on evaluation scores varies with (a) respondent DEI / diversity preferences, (b) whether the respondent's firm targets the same demographic as the applicant, and (c) the triple interaction of respondent DEI priority with firm DEI attention.]

Sign convention: Throughout §3, all displayed coefficients are *flipped* so that positive $\Rightarrow$ blinding raises the relative score for that demographic (mirroring the convention in §1.1). The underlying regressions use Blind as the reference category for `race_gender`; multiplying by -1 converts (nonblind $-$ blind) into (blind $-$ nonblind), i.e. the effect of blinding.

3.1 (2-way Interaction Models) By Individual DEI / Diversity Variables

Takeaway: Mirroring 8A §3's null on preference heterogeneity, most individual DEI / diversity proxies do not meaningfully moderate the effect of blinding on demographic groups. Two exceptions stand out. For HR, respondent-perceived firm DEI attention interacts strongly with the blinding effect on Hispanic Male applicants ($+0.188^{***}$ for Firm Race DEI Attention; $+0.133^{***}$ for Firm Gender DEI Attention) — a firm-level substitution pattern in which HR evaluators at firms with high DEI attention penalize Hispanic Males under nonblind, and blinding mechanically recovers those penalties. For Engineers, individual-level pro-DEI proxies (Race DEI, Gender DEI, Diverse Workplace, Diversity Priority) consistently interact *negatively* with the blinding effect on Hispanic Males (-0.057 to -0.081, $p < 0.05$ or tighter on each) — an individual-level DEI-premium pattern in which pro-DEI Engineers are more generous to Hispanic Males under nonblind. The two channels operate in different places: firm-level substitution for HR; individual-level DEI premium for Engineers.

Reproduction with three DEI / Diversity Index Options. The same regression as above, replacing the six individual DEI proxies with three composite indices: *Index A: Stated* — the four-item average of stated DEI views; *Index B: Stated + Revealed* — equal-weight composite of Index A and the revealed-priority ranking of "Diversity and Representation'' from the goals question; *Index C: PC1* — first principal component of the same five items.

Outcome Var: Assigned Overall Score (Z) — Three DEI / Diversity Index Options, signs flipped so + ⇒ effect of blinding

	Panel A: HR			**Panel B: Engineer**		
	Index A	Index B	Index C	Index A	Index B	Index C
	(1)	(2)	(3)	(4)	(5)	(6)
White Male	0.040	0.039	0.040	0.086***	0.086***	0.086***
	(0.056)	(0.056)	(0.056)	(0.028)	(0.028)	(0.028)
White Female	0.014	0.014	0.014	0.023	0.022	0.023
	(0.051)	(0.051)	(0.051)	(0.036)	(0.035)	(0.035)
Asian Male	0.016	0.015	0.016	-0.017	-0.017	-0.017
	(0.060)	(0.059)	(0.060)	(0.028)	(0.028)	(0.028)
Asian Female	-0.023	-0.023	-0.023	-0.030	-0.030	-0.030
	(0.048)	(0.048)	(0.048)	(0.029)	(0.029)	(0.029)
Hispanic Male	-0.113**	-0.113**	-0.113**	-0.004	-0.004	-0.004
	(0.048)	(0.048)	(0.048)	(0.030)	(0.029)	(0.029)
Hispanic Female	0.091*	0.089*	0.091*	-0.070**	-0.069**	-0.070**
	(0.050)	(0.050)	(0.050)	(0.033)	(0.034)	(0.034)
White Male × DEI Index	0.099	0.089	0.100	0.013	0.002	0.009
	(0.081)	(0.072)	(0.079)	(0.035)	(0.028)	(0.033)
White Female × DEI Index	-0.027	0.007	-0.018	0.025	0.054	0.036
	(0.052)	(0.050)	(0.052)	(0.046)	(0.046)	(0.047)
Asian Male × DEI Index	0.002	-0.018	-0.005	0.009	0.021	0.012
	(0.064)	(0.060)	(0.064)	(0.026)	(0.027)	(0.026)
Asian Female × DEI Index	0.008	0.010	0.008	-0.008	-0.0010	-0.005
	(0.045)	(0.049)	(0.046)	(0.022)	(0.024)	(0.022)
Hispanic Male × DEI Index	0.023	0.014	0.022	-0.064**	-0.088***	-0.073**
	(0.064)	(0.058)	(0.063)	(0.028)	(0.030)	(0.029)
Hispanic Female × DEI Index	-0.134**	-0.122**	-0.135**	0.041	0.029	0.038
	(0.063)	(0.058)	(0.062)	(0.035)	(0.033)	(0.034)
Resume (Content) FE	✓	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓	✓
Respondent Demographic FE	✓	✓	✓	✓	✓	✓
Observations	2,466	2,466	2,466	2,772	2,772	2,772
Adjusted R^2	0.40900	0.40859	0.40898	0.73712	0.73756	0.73726

Index A is the existing four-item average of stated DEI views. Index B averages Index A with the respondent's revealed-priority ranking of "Diversity and Representation'' from the seven-goals question. Index C is the first principal component of the same five items, sign-aligned so higher = more pro-DEI. Coefficients are flipped via I(-sc_overall_z) so positive => effect of blinding. Standard errors clustered by respondent. Sample: both blind and nonblind arms.

Reproduction (URM cut) with three DEI / Diversity Index Options.

Outcome Var: Assigned Overall Score (Z) — URM Cut, Three DEI / Diversity Index Options, signs flipped so + ⇒ effect of blinding

	Panel A: HR			Panel B: Engineer		
	Index A	Index B	Index C	Index A	Index B	Index C
	(1)	(2)	(3)	(4)	(5)	(6)
URM (WF, AF, HM, HF)	-0.013	-0.013	-0.013	-0.019*	-0.019*	-0.019*
	(0.019)	(0.019)	(0.019)	(0.010)	(0.010)	(0.010)
Non-URM	0.026	0.026	0.026	0.035*	0.035*	0.035*
	(0.037)	(0.037)	(0.037)	(0.019)	(0.018)	(0.019)
URM × DEI Index	-0.027	-0.019	-0.025	-0.006	-0.006	-0.005
	(0.023)	(0.022)	(0.023)	(0.013)	(0.010)	(0.012)
Non-URM × DEI Index	0.052	0.037	0.049	0.011	0.011	0.010
	(0.043)	(0.042)	(0.043)	(0.024)	(0.019)	(0.022)
Resume (Content) FE	✓	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓	✓
Respondent Demographic FE	✓	✓	✓	✓	✓	✓
Observations	2,466	2,466	2,466	2,772	2,772	2,772
Adjusted R^2	0.40775	0.40741	0.40767	0.73661	0.73661	0.73661

URM-Tech = {White Female, Asian Female, Hispanic Male, Hispanic Female}. Sample: both blind and nonblind arms. Coefficients flipped via I(-sc_overall_z) so positive => effect of blinding. Standard errors clustered by respondent.

Visual: URM–Non-URM gap under nonblind by DEI Index (Indices A, B, C). Each row is one of the three DEI indices; columns are role. The single curve in each panel plots the URM–Non-URM premium under nonblind — exactly the gap that blinding eliminates by random assignment. Positive values: URM rated higher than Non-URM under nonblind. Negative values: Non-URM rated higher. Shaded band is the 95% CI; the dashed line at zero marks where blinding has no reshuffling effect.

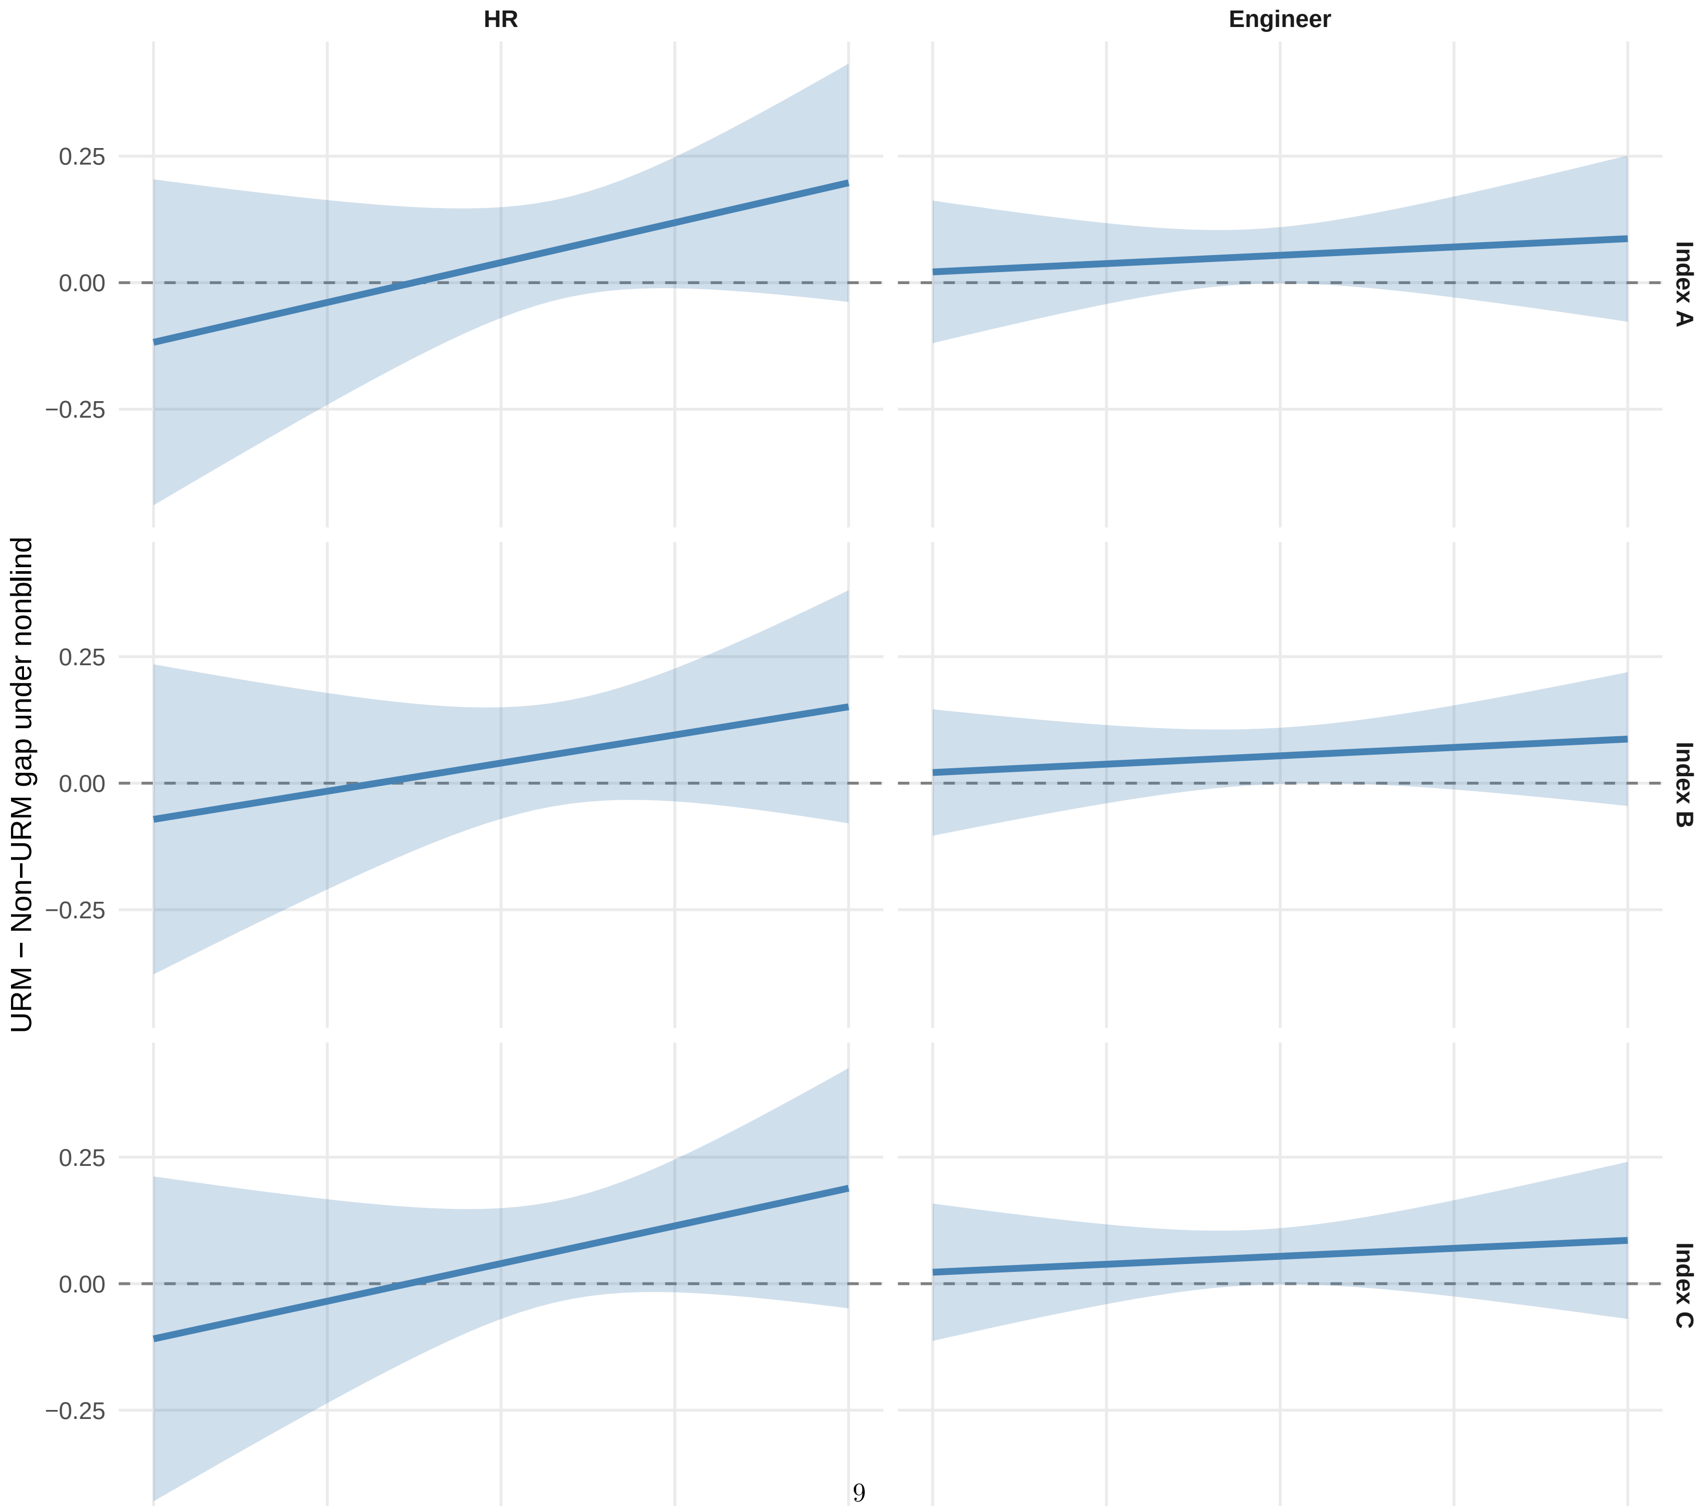
HR
Engineer
Index A
Index B
Index C
URM − Non−URM gap under nonblind
0.25
0.00
−0.25

3.2 (2-way Interaction Models) By Whether the Firm Targets the Same Applicant Demographic

Takeaway: The substitution hypothesis is a *null* at the level it can be cleanly tested. The bottom table's URM-aggregate match coefficient (URM × Firm Targets URM) is insignificant in both roles, and so is the Non-URM spillover. The top table is descriptive only: per-demographic match cells are identified off too few respondents to support inference (we restrict it to the URM-aligned dimensions Female and Hispanic and drop the rest, since White / Asian / Male match cells are either non-URM-aligned or fire on ≤ 2 unique respondents per role). The one mover — Engineer's `applicant female` × `firm targets female` — points in the substitution direction but is identified off ≈ 5 respondents and should be read as suggestive. The cleaner Engineer signal lives in the *main effects*: blinding *lowers* URM applicants' Overall scores and *raises* Non-URM scores (significant at the 5% level), reproducing the Engineer non-concordance premium documented in §1.1 and §2 net of firm-targeting controls.

Firm Targets × Same Applicant Demographic — signs flipped so + ⇒ effect of blinding

	Panel A: HR	Panel B: Engineer
	(1)	(2)
White Male	0.039	0.085***
	(0.058)	(0.028)
White Female	0.022	0.010
	(0.053)	(0.035)
Asian Male	0.016	-0.018
	(0.060)	(0.028)
Asian Female	-0.014	-0.042
	(0.048)	(0.030)
Hispanic Male	-0.119**	-0.003
	(0.048)	(0.031)
Hispanic Female	0.090	-0.081**
	(0.058)	(0.032)
Applicant Female × Firm Targets Female	-0.079	0.186**
	(0.145)	(0.080)
Applicant Hispanic × Firm Targets Hispanic	0.032	-0.003
	(0.102)	(0.087)
Resume (Content) FE	✓	✓
Display Position FE	✓	✓
Respondent Demographic FE	✓	✓
Observations	2,466	2,772
Adjusted R^2	0.40790	0.73747

Each Applicant X x Firm Targets X is a match interaction, identified off applicant rows whose demographic equals the firm's targeting dimension. Standard errors clustered by respondent. Sample: both blind and nonblind arms.

URM-Tech Binary Cut × Firm Targets (Female OR Hispanic) — signs flipped

	Panel A: HR	Panel B: Engineer
	(1)	(2)
Non-URM	0.005	0.042**
	(0.040)	(0.019)
URM (WF, AF, HM, HF)	-0.003	-0.022**
	(0.020)	(0.010)
Non-URM × Firm Targets URM (spillover)	0.098	-0.096
	(0.105)	(0.069)
URM × Firm Targets URM	-0.051	0.050
	(0.054)	(0.035)
Resume (Content) FE	✓	✓
Display Position FE	✓	✓
Respondent Demographic FE	✓	✓
Observations	2,466	2,772
Adjusted R^2	0.40748	0.73673

URM = {WF, AF, HM, HF}. "Firm Targets URM' ' => firm_targets_female = 1 OR firm_targets_hispanic = 1. The URM row is the aggregate match; the Non-URM row is a spillover test (does targeting URM groups also shift the blinding effect on Non-URM applicants?). Standard errors clustered by respondent. Sample: both blind and nonblind arms.

3.3 (Triple Interaction Models) By Respondent's Priority for DEI $\times$ Firm Attention on DEI

Takeaway: For HR, the alignment of respondent's DEI priority with firm's DEI attention substitutes for blinding: at pro-DEI individuals $\times$ pro-DEI firms, blinding has near-zero effect on either URM or Non-URM applicants, while at anti-DEI $\times$ anti-DEI it flips a ~ 0.40 SD nonblind premium favoring Non-URM (URM $\times$ DEI Index $\times$ Firm DEI Attn = $+0.09^{***}$; Non-URM mirror = -0.19^{***}). Engineers show null moderation on both groups — DEI-driven heterogeneity in treatment effects lives in HR's evaluations, not Engineers'.

Hypothesis: Same **Substitution vs. Backlash** structure as 8A §3, now applied to the blinding effect on URM applicants.

- **Substitution**: When the firm already prioritizes DEI, respondents who *also* personally prioritize DEI absorb the firm's push and ease off their own URM premium under nonblind. Blinding therefore bites *less* for them (small effect of blinding on URM), because their nonblind ratings were already closer to the demographically blind baseline. *Expect a negative sign on the triple interaction term.*
- **Backlash**: When the firm prioritizes DEI, respondents who do *not* personally prioritize DEI react by marking URM applicants lower under nonblind. Blinding therefore bites *more* for them (large positive effect of blinding on URM). *Expect a positive sign on the triple interaction term.*

Reproduction with three DEI / Diversity Index Options. The same triple-interaction structure as above, but replacing the single revealed-priority variable (`div_priority_z`) with each of three composite DEI indices in turn. One sub-table per index; each reports HR vs. Eng under the combined firm-DEI-attention-level proxy (`firm_dei_attention_z`).

Index A — signs flipped so + ⇒ effect of blinding

	Firm DEI Attn Lvl	
	HR	Eng
	(1)	(2)
URM (WF, AF, HM, HF)	-0.042**	-0.007
	(0.021)	(0.010)
Non-URM	0.084**	0.014
	(0.040)	(0.019)
URM × DEI Index	-0.037*	-0.021
	(0.019)	(0.019)
Non-URM × DEI Index	0.069*	0.040
	(0.035)	(0.036)
URM × Firm DEI Attn	-0.045	0.010
	(0.030)	(0.012)
Non-URM × Firm DEI Attn	0.090	-0.018
	(0.061)	(0.024)
URM × DEI Index × Firm DEI Attn	0.094***	-0.009
	(0.031)	(0.014)
Non-URM × DEI Index × Firm DEI Attn	-0.188***	0.017
	(0.061)	(0.025)
Resume (Content) FE	✓	✓
Display Position FE	✓	✓
Respondent Demographic FE	✓	✓
Observations	2,160	2,358
Adjusted R^2	0.40224	0.73463

Firm DEI Attn proxy is firm_dei_attention_z (within-cell z-score of the equal-weight average of firm_attention_race_z and firm_attention_gender_z). Coefficients flipped via I(-sc_overall_z) so positive => effect of blinding. Standard errors clustered by respondent. Sample: both arms.

Index B — signs flipped so + ⇒ effect of blinding

	Firm DEI Attn Lvl	
	HR	Eng
	(1)	(2)
URM (WF, AF, HM, HF)	-0.034*	-0.014
	(0.020)	(0.011)
Non-URM	0.067*	0.027
	(0.038)	(0.022)
URM × DEI Index	-0.037*	-0.015
	(0.021)	(0.012)
Non-URM × DEI Index	0.070*	0.028
	(0.039)	(0.023)
URM × Firm DEI Attn	-0.012	0.005
	(0.020)	(0.014)
Non-URM × Firm DEI Attn	0.024	-0.009
	(0.041)	(0.027)
URM × DEI Index × Firm DEI Attn	0.074***	0.004
	(0.022)	(0.010)
Non-URM × DEI Index × Firm DEI Attn	-0.147***	-0.006
	(0.044)	(0.019)
Resume (Content) FE	✓	✓
Display Position FE	✓	✓
Respondent Demographic FE	✓	✓
Observations	2,160	2,358
Adjusted R^2	0.40074	0.73459

Firm DEI Attn proxy is firm_dei_attention_z (within-cell z-score of the equal-weight average of firm_attention_race_z and firm_attention_gender_z). Coefficients flipped via I(-sc_overall_z) so positive => effect of blinding. Standard errors clustered by respondent. Sample: both arms.

Index C — signs flipped so + ⇒ effect of blinding		
	Firm DEI Attn Lvl	
	HR	Eng
	(1)	(2)
URM (WF, AF, HM, HF)	-0.044**	-0.008
	(0.021)	(0.010)
Non-URM	0.087**	0.015
	(0.039)	(0.019)
URM × DEI Index	-0.037*	-0.021
	(0.019)	(0.018)
Non-URM × DEI Index	0.070*	0.040
	(0.036)	(0.033)
URM × Firm DEI Attn	-0.040	0.010
	(0.028)	(0.012)
Non-URM × Firm DEI Attn	0.081	-0.017
	(0.055)	(0.024)
URM × DEI Index × Firm DEI Attn	0.096***	-0.008
	(0.029)	(0.012)
Non-URM × DEI Index × Firm DEI Attn	-0.192***	0.015
	(0.058)	(0.023)
Resume (Content) FE	✓	✓
Display Position FE	✓	✓
Respondent Demographic FE	✓	✓
Observations	2,160	2,358
Adjusted R^2	0.40243	0.73462

Firm DEI Attn proxy is firm_dei_attention_z (within-cell z-score of the equal-weight average of firm_attention_race_z and firm_attention_gender_z). Coefficients flipped via I(-sc_overall_z) so positive => effect of blinding. Standard errors clustered by respondent. Sample: both arms.

Visual: URM–Non-URM gap under nonblind by DEI Index × Firm DEI Attn. One figure per DEI index (A, B, C). Each panel plots the URM–Non-URM premium under nonblind = the gap blinding eliminates. Rows: Firm DEI Attn fixed at +1 SD (high, top) and −1 SD (low, bottom). Columns: role. Substitution shows up as the HR gap curve flattening toward zero in the top-right region (high firm DEI attn, high individual DEI Index).

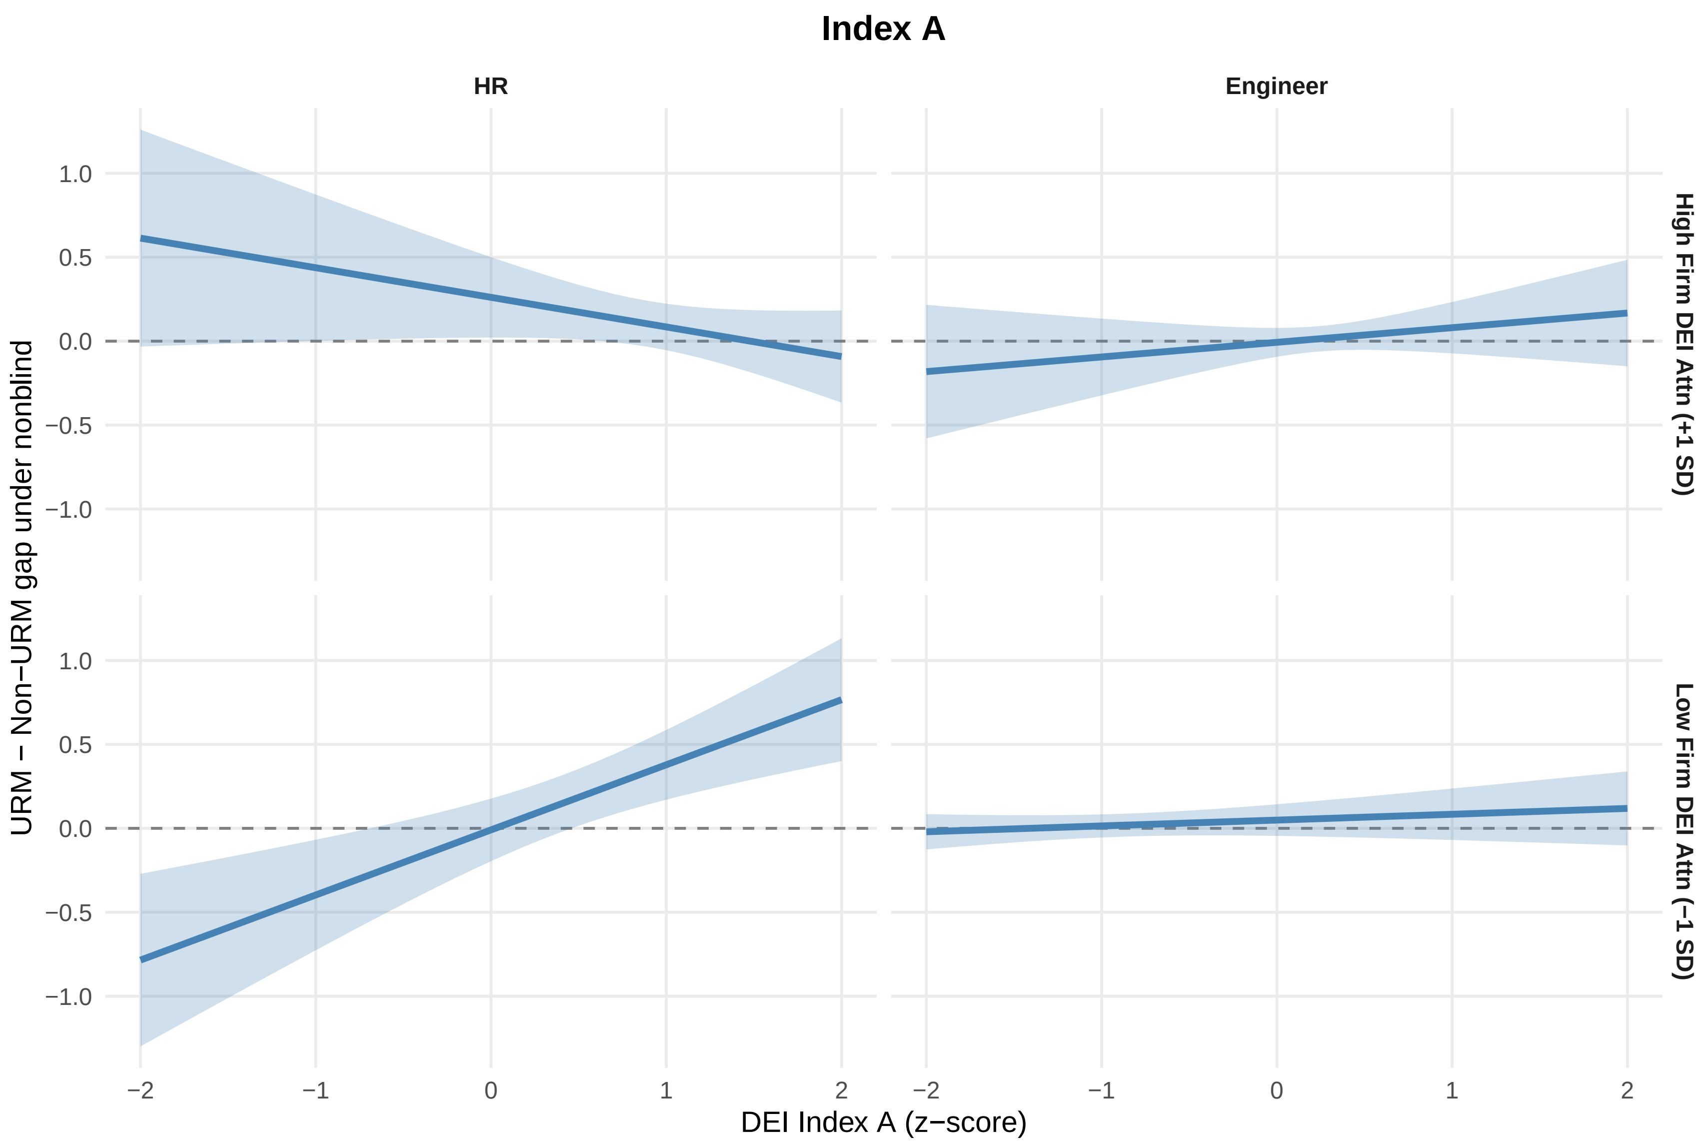
Index A
HR
Engineer
High Firm DEI Attn (+1 SD)
Low Firm DEI Attn (−1 SD)
URM − Non−URM gap under nonblind
DEI Index A (z−score)
1.0
0.5
0.0
−0.5
−1.0
−2
−1
0
1
2

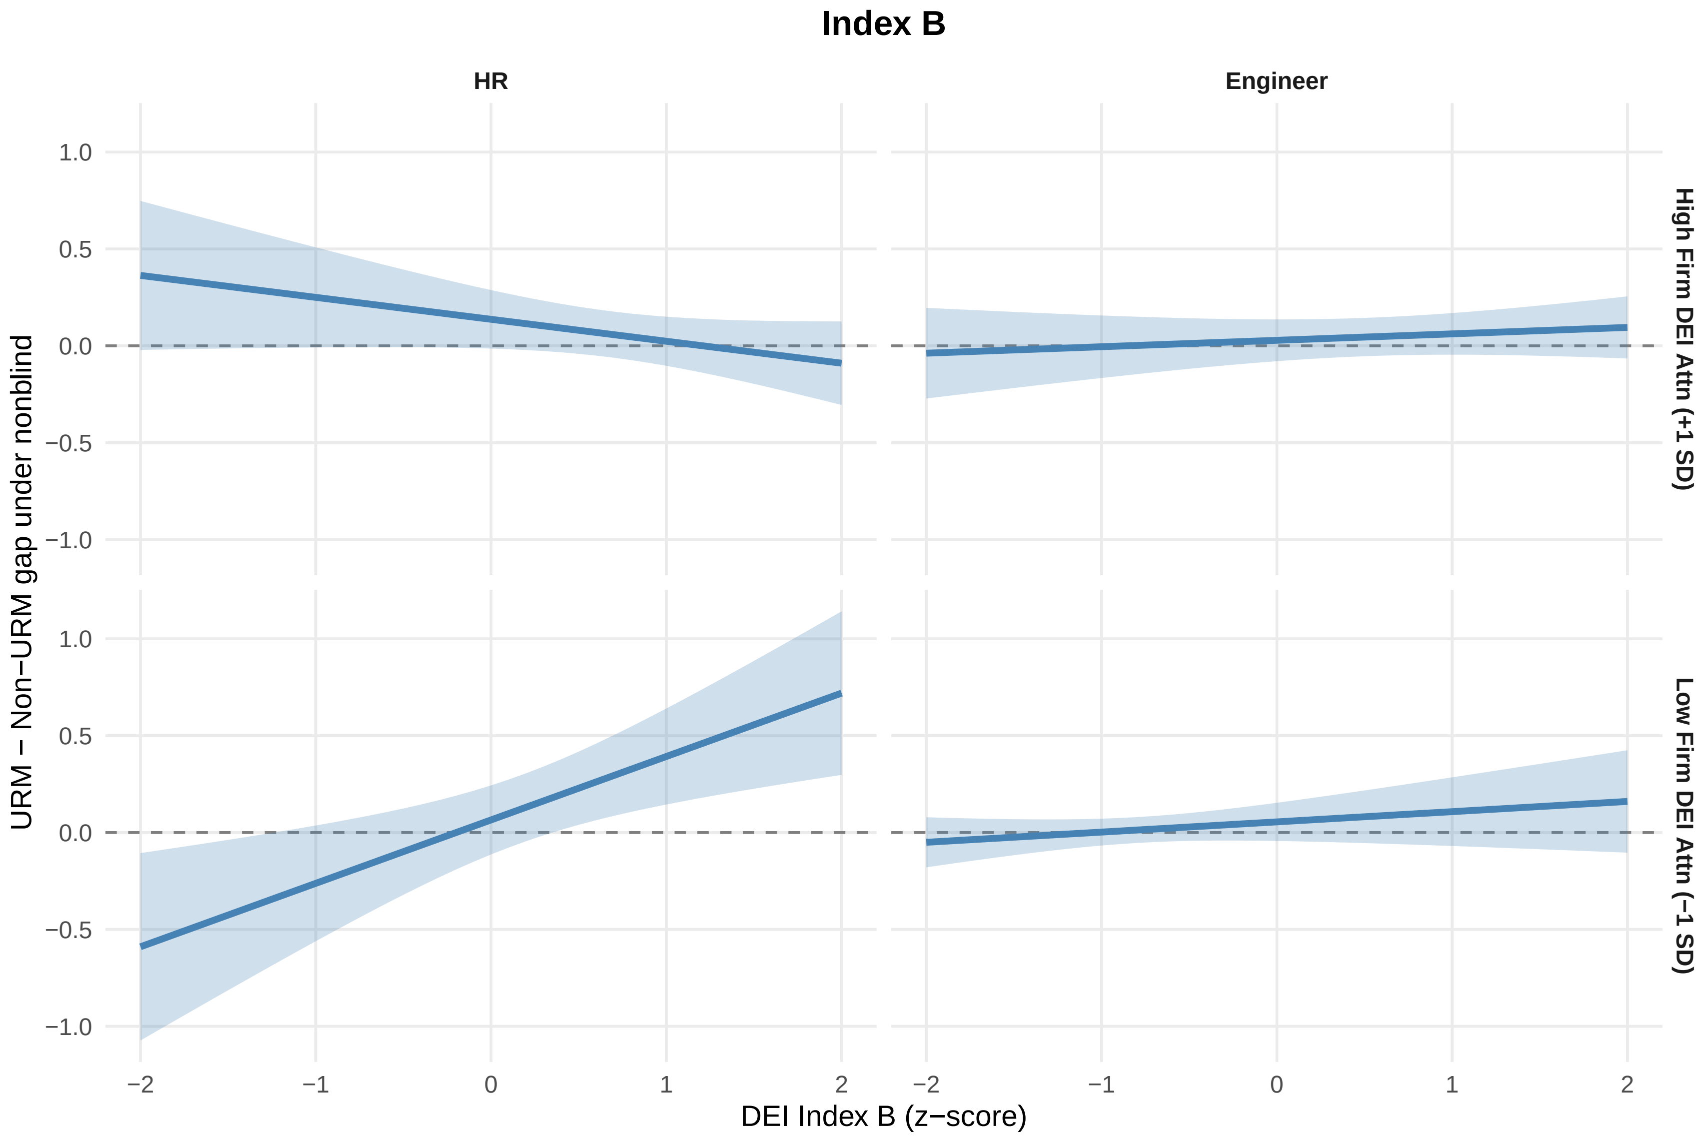
Index B
HR
Engineer
High Firm DEI Attn (+1 SD)
Low Firm DEI Attn (−1 SD)
URM − Non−URM gap under nonblind
DEI Index B (z−score)
1.0
0.5
0.0
−0.5
−1.0
−2
−1
0
1
2

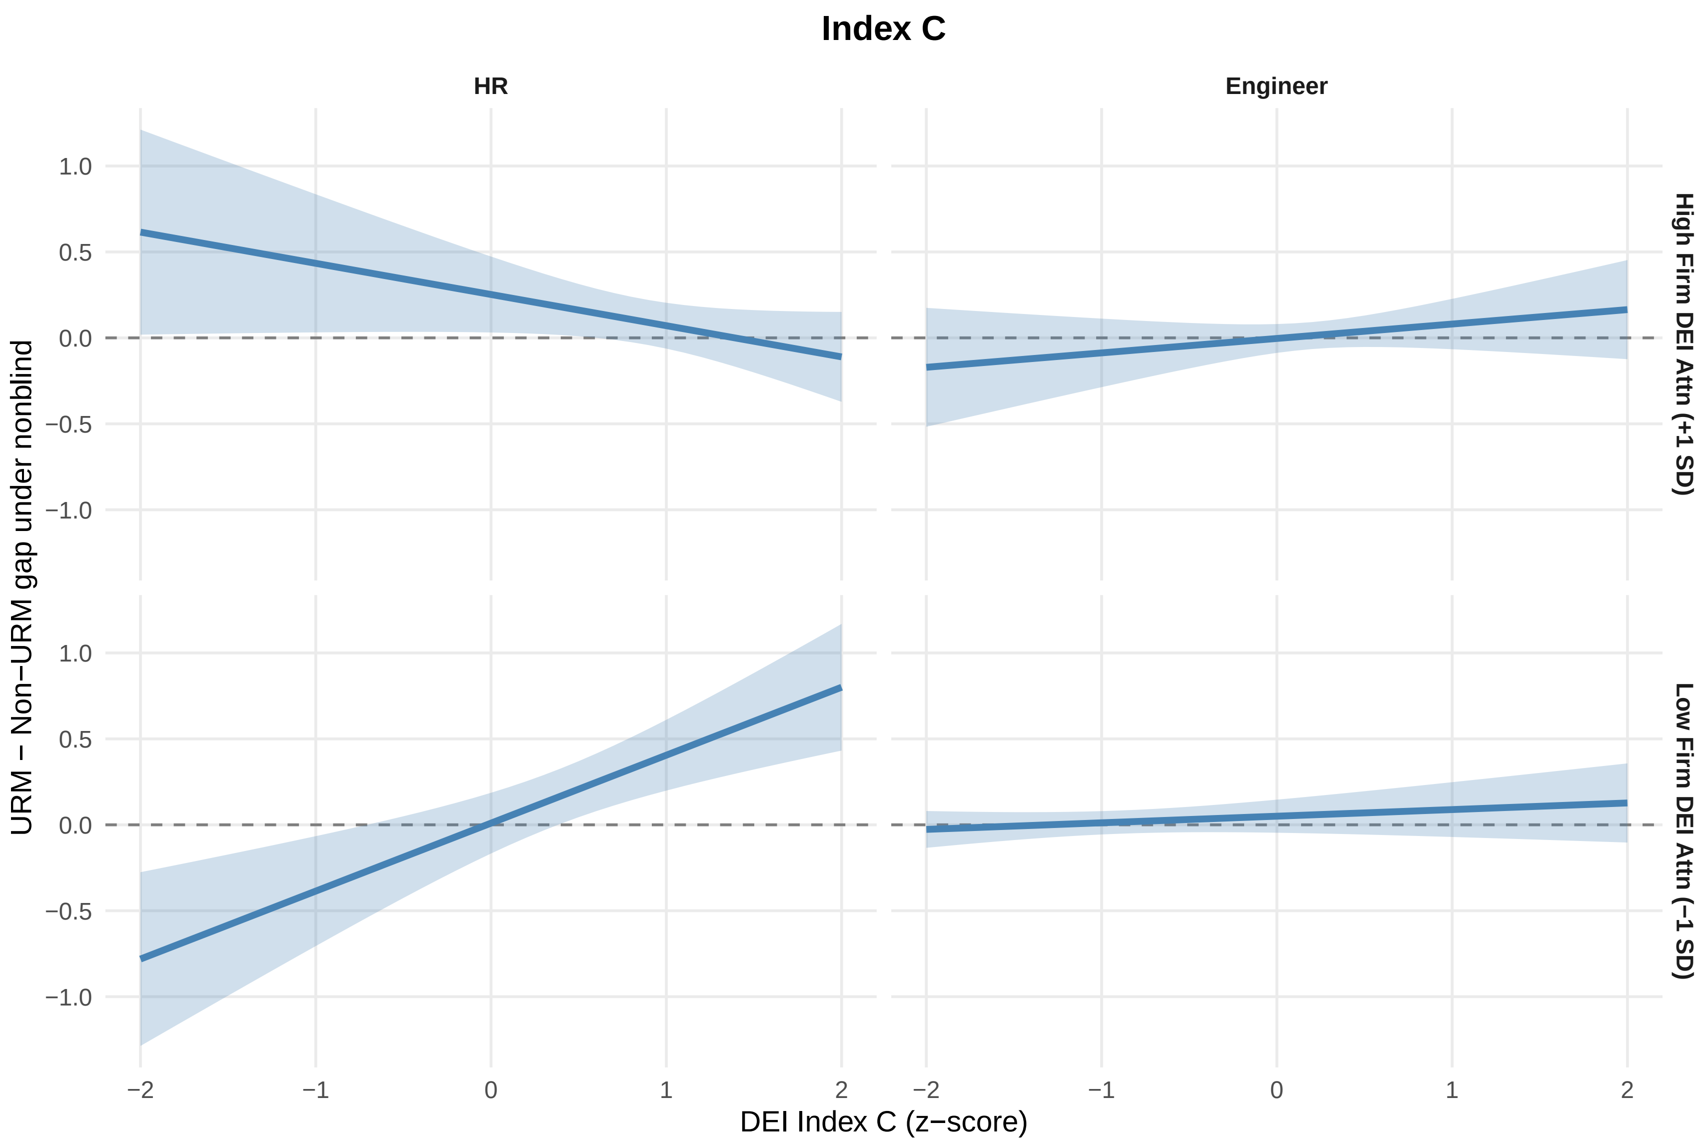

Index C
HR
Engineer
High Firm DEI Attn (+1 SD)
Low Firm DEI Attn (−1 SD)
URM − Non−URM gap under nonblind
DEI Index C (z−score)
1.0
0.5
0.0
−0.5
−1.0
−2
−1
0
1
2

4 Treatment Effect on Productivity of Hired