Stage-by-Stage Impact of Blind Hiring

1 Impact of Blind Hiring on Evaluations

1.1 On All Dimensions (Z)

Takeaway: Our experiment demonstrates that withholding demographic information *reshuffles* candidates' relative evaluations across stages. In the HR resume-screening stage, Hispanic Male candidates are the largest losers: under blinding, their overall evaluation scores fall by 11.8% of a standard deviation relative to when their identity is revealed (significant at the 5% level). In the Engineer technical evaluation stage, the largest gainers are White Male candidates, whose scores rise by 8.5% of a SD under blinding (significant at the 1% level), while most other demographic groups lose ground. A standard caveat applies: because hiring is often an ordinal ranking problem — where one candidate's gain necessarily corresponds to another candidate's lost opportunity — the absence of statistical significance on a given row does not mean the corresponding group is unaffected by blinding.

Panel A: HR

	Coding (Z)	Fit (Z)	Stay (Z)	Overall (Z)	Advance (0/1)
White Male	0.039	0.075	0.072	0.040	−0.005
	(0.056)	(0.057)	(0.059)	(0.057)	(0.050)
White Female	0.020	0.011	0.025	0.012	0.002
	(0.046)	(0.052)	(0.065)	(0.051)	(0.049)
Asian Male	0.034	0.031	−0.049	0.013	−0.021
	(0.051)	(0.061)	(0.061)	(0.059)	(0.048)
Asian Female	0.018	−0.007	0.027	−0.020	−0.017
	(0.052)	(0.055)	(0.061)	(0.048)	(0.052)
Hispanic Male	−0.142***	−0.064	−0.087	−0.118**	−0.049
	(0.050)	(0.053)	(0.062)	(0.049)	(0.048)
Hispanic Female	0.040	−0.050	0.013	0.087*	−0.001
	(0.058)	(0.066)	(0.074)	(0.052)	(0.055)
Observations	2466	2430	2394	2466	2466
R2	0.397	0.310	0.194	0.421	0.226
Resume (Content) FE	✓	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓	✓
Respondent FE	×	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓	✓

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Panel B: Engineer

	Coding (Z)	Fit (Z)	Stay (Z)	Overall (Z)
White Male	0.070**	0.041	0.051	0.085***
	(0.029)	(0.041)	(0.049)	(0.028)
White Female	0.027	0.043	0.015	0.019
	(0.031)	(0.044)	(0.052)	(0.036)
Asian Male	−0.069**	0.077*	0.000	−0.016
	(0.028)	(0.047)	(0.050)	(0.028)
Asian Female	0.006	−0.094**	−0.055	−0.027
	(0.029)	(0.043)	(0.048)	(0.029)
Hispanic Male	0.009	0.023	0.039	−0.003
	(0.027)	(0.037)	(0.054)	(0.030)
Hispanic Female	−0.053**	−0.097**	−0.060	−0.068**
	(0.026)	(0.041)	(0.058)	(0.034)
Observations	2772	2772	2736	2772
R2	0.787	0.569	0.343	0.742
Resume (Content) FE	✓	✓	✓	✓
Display Position FE	✓	✓	✓	✓
Respondent FE	×	×	×	×
Respondent Characteristics	✓	✓	✓	✓

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Standard errors clustered by respondent.

Note on Interpreting Coefficients: The underlying regression uses Blind resumes as the reference category, so raw coefficients capture (nonblind − blind) or the effect of being blinded for each demographic on relative overall evaluation scores. We then multiply by −1 to report (blind − nonblind), aligning the sign convention with the standard "effect of blinding" framing: positive ⇒ blinding raises the relative score for that demographic. "Respondent Characteristics" comprises age (with polynomial), gender, race, education, and income.

1.2 Does Blinding Lead to More Accurate Coding Evaluations?

Takeaway: Evaluators do extract signal about candidates' actual coding ability — Engineers much more strongly than HR, consistent with Engineers seeing the code while HR sees only the resume. Blinding, however, does *not* meaningfully change this signal-extraction ability in either role: the `true_coding_z` × `blind` interaction is small, negative, and statistically insignificant for both HR and Engineers. If blinding — which in our experiment amounts to removing random, uninformative demographic noise — improved signal-extraction ability, we would expect a positive and statistically significant interaction term. In short, blinding *reshuffles* who receives high scores (Section 1.1) but does not change *how well* those scores track actual coding skill.

Specification: For each role, we fit

$$\text{Hiring Evaluation Score (Z)} \sim \text{True Coding Ability (Z)} \times \not\Vdash[\text{Blind}] \mid \text{Respondent FE} + \text{Display Position FE}.$$

The main slope on True Coding Ability (Z) is the evaluator's signal sensitivity in the non-blind arm; the interaction is the shift in sensitivity under blinding. Resume-content FE is deliberately absent, because True Coding Ability (Z) is constant within a resume (one lab-measured score per resume) and a resume FE would absorb the RHS entirely.

	HR	Engineer
True Coding Score (Z)	0.379***	0.886***
	(0.029)	(0.012)
True Coding Score (Z) × Blind	−0.021	−0.031
	(0.037)	(0.020)
Observations	2466	2772
R2	0.139	0.758
Resume (Content) FE	×	×
Display Position FE	✓	✓
Respondent FE	✓	✓
Respondent Characteristics	×	×

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Standard errors clustered by respondent.

Note on Interpreting Coefficients: The main slope on "True Coding Score (Z)" is the evaluator's signal sensitivity in the non-blind arm — a 1 SD increase in actual coding ability maps into a $\hat{\beta}$ SD increase in the assigned coding rating. The interaction row is the shift in that sensitivity under blinding: positive ⇒ blinding sharpens rating-to-truth tracking, negative ⇒ blinding dulls it. The "Resume (Content) FE" row is marked as not included because `overall_score` is a resume-level property (one lab score per resume), so a resume-content FE would absorb the RHS and leave the main slope unidentified. "Respondent Characteristics" is marked as not included because respondent-level variables (age, gender, race, education, income) are all perfectly collinear with the respondent FE and therefore absorbed by it.

2 Stage-by-Stage Impact of Blind Hiring on Coding Outcomes

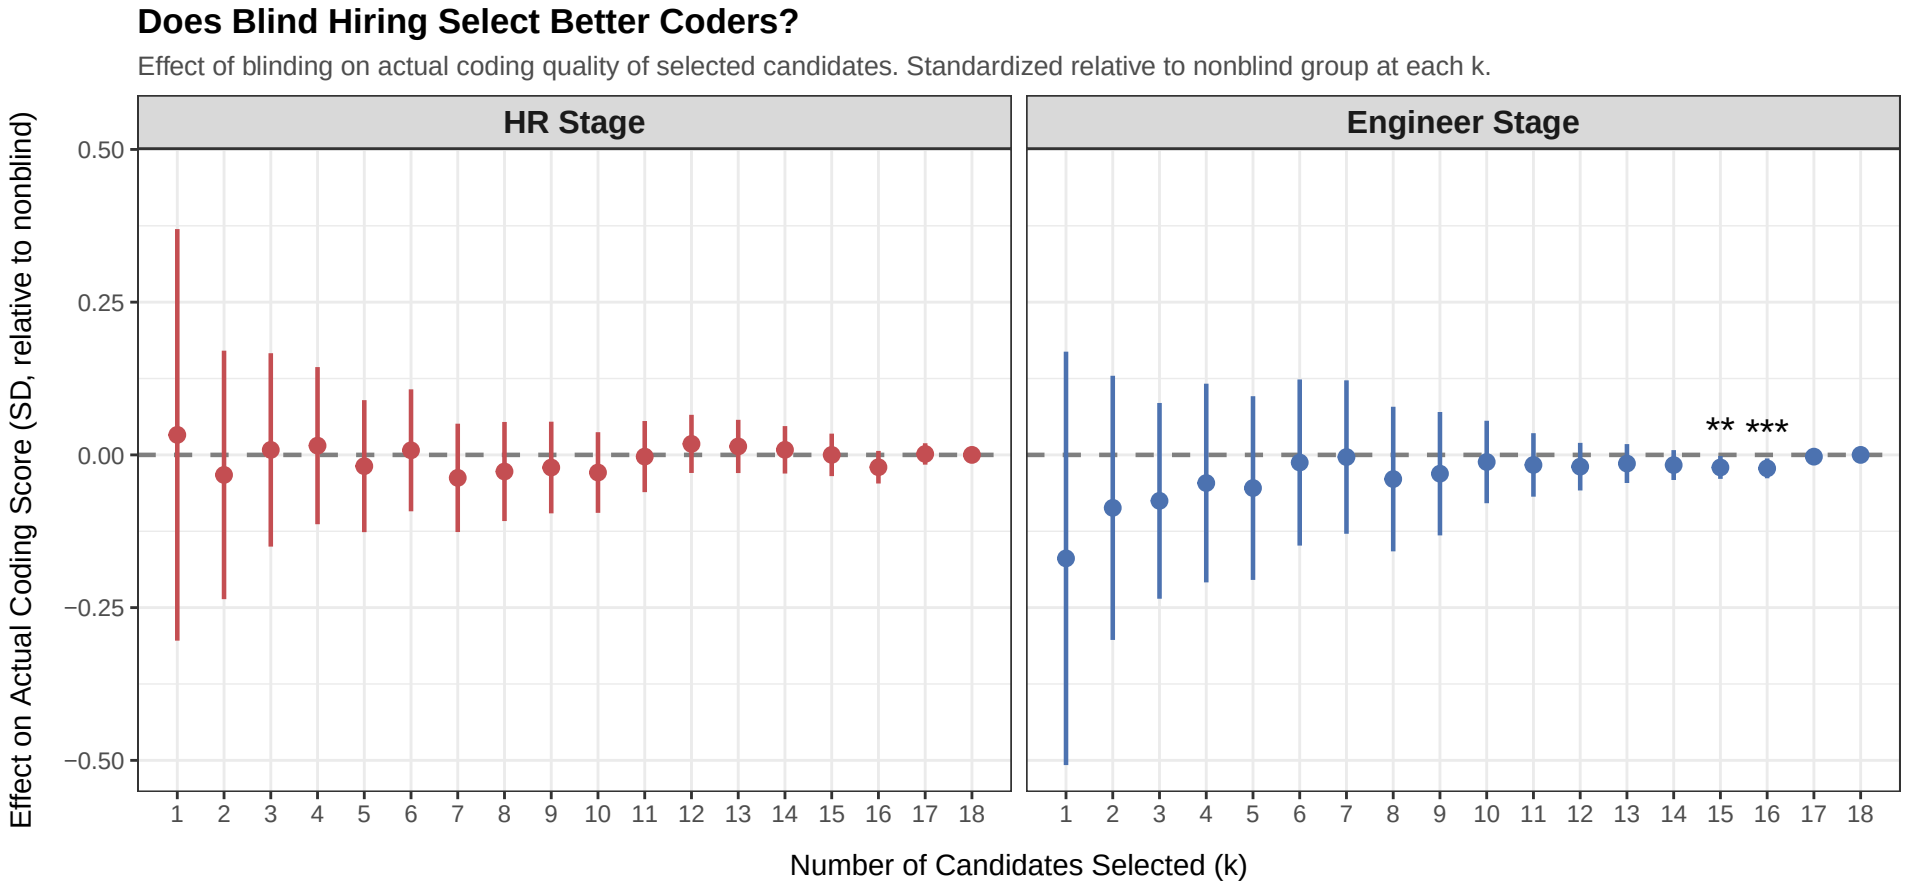

Positive = blind hiring selects candidates with higher actual coding scores.
Significance: * p<0.10, ** p<0.05, *** p<0.01